

Sapienza University of Rome

Master in Artificial Intelligence and Robotics  
Master in Engineering in Computer Science

## Machine Learning

A.Y. 2021/2022

Prof. Luca Iocchi

## 7. Linear models for classification

Luca Iocchi

# Overview

- Linearly separable data
- Linear models
- Least squares
- Perceptron
- Fisher's linear discriminant
- Support Vector Machines

## References

C. Bishop. Pattern Recognition and Machine Learning. Sect. 4.1, 7.1

T. Mitchell. Machine Learning. Section 4.4

## Linear Models for Classification

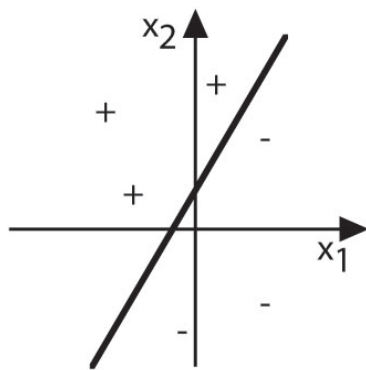
Learning a function  $f : X \rightarrow Y$ , with ...

- $X \subseteq \mathbb{R}^d$
- $Y = \{C_1, \dots, C_k\}$

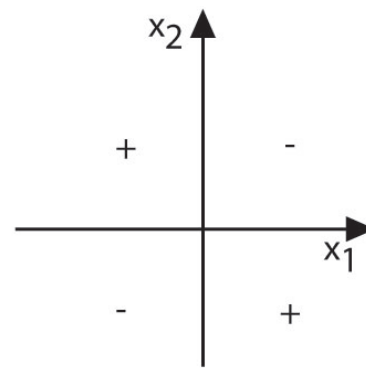
assuming *linearly separable* data.

## Linearly separable data

Instances in a data set are *linearly separable* iff there exists a hyperplane that separates the instance space into two regions, such that differently classified instances are separated



(a)



(b)

## Linear discriminant functions

Linear discriminant function

$$y : X \rightarrow \{C_1, \dots, C_K\}$$

Two classes:

$$y(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0$$

$K$ -class:

$$y_1(\mathbf{x}) = \mathbf{w}_1^T \mathbf{x} + w_{10}$$

...

$$y_K(\mathbf{x}) = \mathbf{w}_K^T \mathbf{x} + w_{K0}$$

## Compact notation

Two classes:

$$y(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0 = \tilde{\mathbf{w}}^T \tilde{\mathbf{x}}, \text{ with:}$$

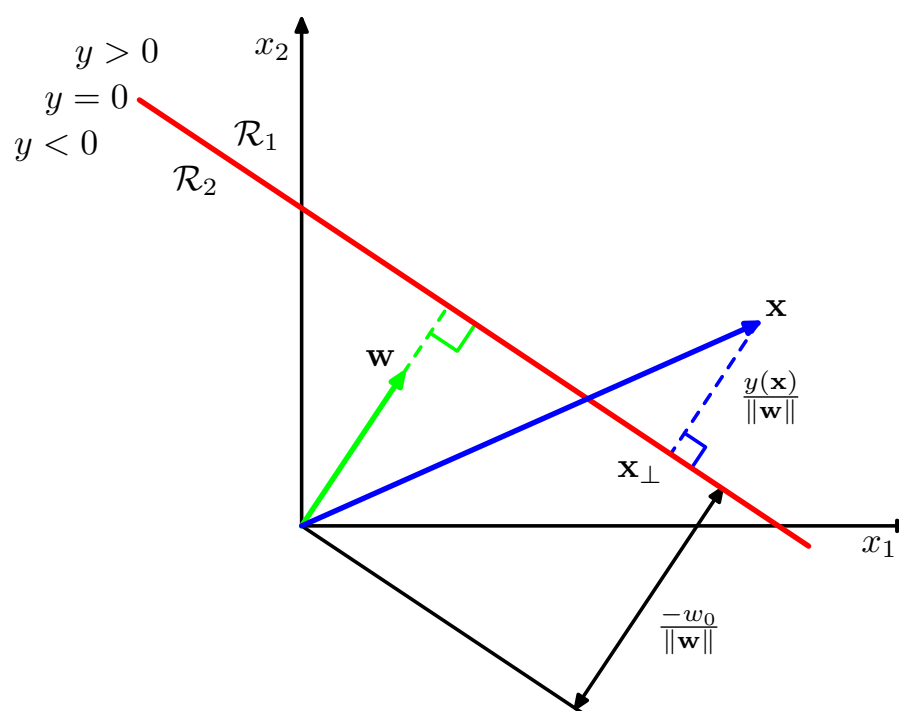
$$\tilde{\mathbf{w}} = \begin{pmatrix} w_0 \\ \mathbf{w} \end{pmatrix}, \tilde{\mathbf{x}} = \begin{pmatrix} 1 \\ \mathbf{x} \end{pmatrix}$$

$K$  classes:

$$\mathbf{y}(\mathbf{x}) = \begin{pmatrix} y_1(\mathbf{x}) \\ \vdots \\ y_K(\mathbf{x}) \end{pmatrix} = \begin{pmatrix} \mathbf{w}_1^T \mathbf{x} + w_{10} \\ \vdots \\ \mathbf{w}_K^T \mathbf{x} + w_{K0} \end{pmatrix} = \begin{pmatrix} \tilde{\mathbf{w}}_1^T \\ \vdots \\ \tilde{\mathbf{w}}_K^T \end{pmatrix} \tilde{\mathbf{x}} = \tilde{\mathbf{W}}^T \tilde{\mathbf{x}}, \text{ with:}$$

$$\tilde{\mathbf{W}}^T = \begin{pmatrix} \tilde{\mathbf{w}}_1^T \\ \vdots \\ \tilde{\mathbf{w}}_K^T \end{pmatrix}, \text{ i.e.: } \tilde{\mathbf{W}} = (\tilde{\mathbf{w}}_1, \dots, \tilde{\mathbf{w}}_K)$$

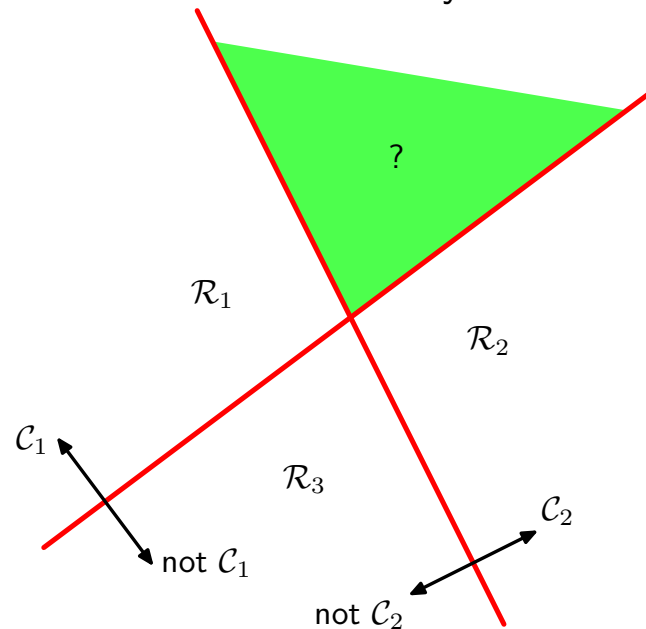
## Linear discriminant functions



## Multiple classes

Cannot use combinations of binary linear models.

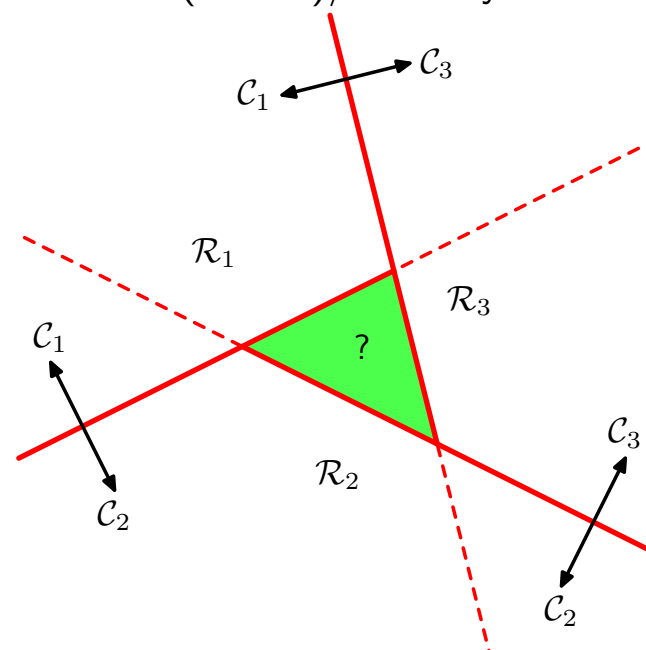
*One-versus-the-rest* classifiers:  $K - 1$  binary classifiers:  $C_k$  vs. not- $C_k$



## Multiple classes

Cannot use combinations of binary linear models.

*One-versus-one* classifiers:  $K(K - 1)/2$  binary classifiers:  $C_k$  vs.  $C_j$



## Multiple classes

$K$ -class discriminant comprising  $K$  linear functions ( $\mathbf{x}$  not in dataset)

$$\mathbf{y}(\mathbf{x}) = \begin{pmatrix} y_1(\mathbf{x}) \\ \vdots \\ y_K(\mathbf{x}) \end{pmatrix} = \begin{pmatrix} \tilde{\mathbf{w}}_1^T \tilde{\mathbf{x}} \\ \vdots \\ \tilde{\mathbf{w}}_K^T \tilde{\mathbf{x}} \end{pmatrix} = \tilde{\mathbf{W}}^T \tilde{\mathbf{x}}$$

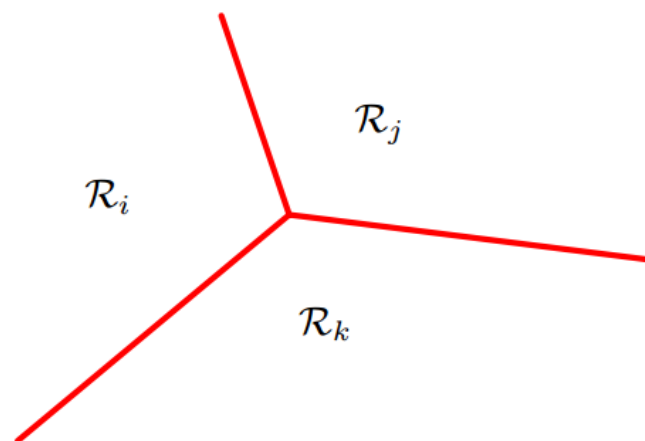
Classify  $\mathbf{x}$  as  $C_k$  if  $y_k(\mathbf{x}) > y_j(\mathbf{x})$  for all  $j \neq k$  ( $j, k = 1, \dots, K$ )

Decision boundary between  $C_k$  and  $C_j$  (hyperplane in  $\Re^{D-1}$ ):

$$(\tilde{\mathbf{w}}_k - \tilde{\mathbf{w}}_j)^T \tilde{\mathbf{x}} = 0$$

## Multiple classes

Example of  $K$ -class discriminant



# Learning linear discriminants

Given a multi-class classification problem and data set  $D$  with linearly separable data,

determine  $\tilde{\mathbf{W}}$  such that  $\mathbf{y}(\mathbf{x}) = \tilde{\mathbf{W}}^T \tilde{\mathbf{x}}$  is the  $K$ -class discriminant.

## Approaches to learn linear discriminants

- Least squares
- Perceptron
- Fisher's linear discriminant
- Support Vector Machines

## Least squares

Given  $D = \{(\mathbf{x}_n, \mathbf{t}_n)_{n=1}^N\}$ , find the linear discriminant

$$\mathbf{y}(\mathbf{x}) = \tilde{\mathbf{W}}^T \tilde{\mathbf{x}}$$

1-of-K coding scheme for  $\mathbf{t}$ :  $\mathbf{x} \in C_k \rightarrow t_k = 1, t_j = 0$  for all  $j \neq k$ .  
E.g.,  $\mathbf{t}_n = (0, \dots, 1, \dots, 0)^T$

$$\tilde{\mathbf{X}} = \begin{pmatrix} \tilde{\mathbf{x}}_1^T \\ \dots \\ \tilde{\mathbf{x}}_N^T \end{pmatrix} \quad \mathbf{T} = \begin{pmatrix} \mathbf{t}_1^T \\ \dots \\ \mathbf{t}_N^T \end{pmatrix}$$

## Least squares

Minimize sum-of-squares error function

$$E(\tilde{\mathbf{W}}) = \frac{1}{2} \text{Tr} \left\{ (\tilde{\mathbf{X}}\tilde{\mathbf{W}} - \mathbf{T})^T (\tilde{\mathbf{X}}\tilde{\mathbf{W}} - \mathbf{T}) \right\}$$

Closed-form solution:

$$\tilde{\mathbf{W}} = \underbrace{(\tilde{\mathbf{X}}^T \tilde{\mathbf{X}})^{-1} \tilde{\mathbf{X}}^T}_{\tilde{\mathbf{X}}^\dagger} \mathbf{T}$$

$$\mathbf{y}(\mathbf{X}) = \tilde{\mathbf{W}}^T \tilde{\mathbf{X}} = \mathbf{T}^T (\tilde{\mathbf{X}}^\dagger)^T \tilde{\mathbf{X}}$$



# Least squares

Classification of new instance  $\mathbf{x}$  not in dataset:

Use learnt  $\tilde{\mathbf{W}}$  to compute:

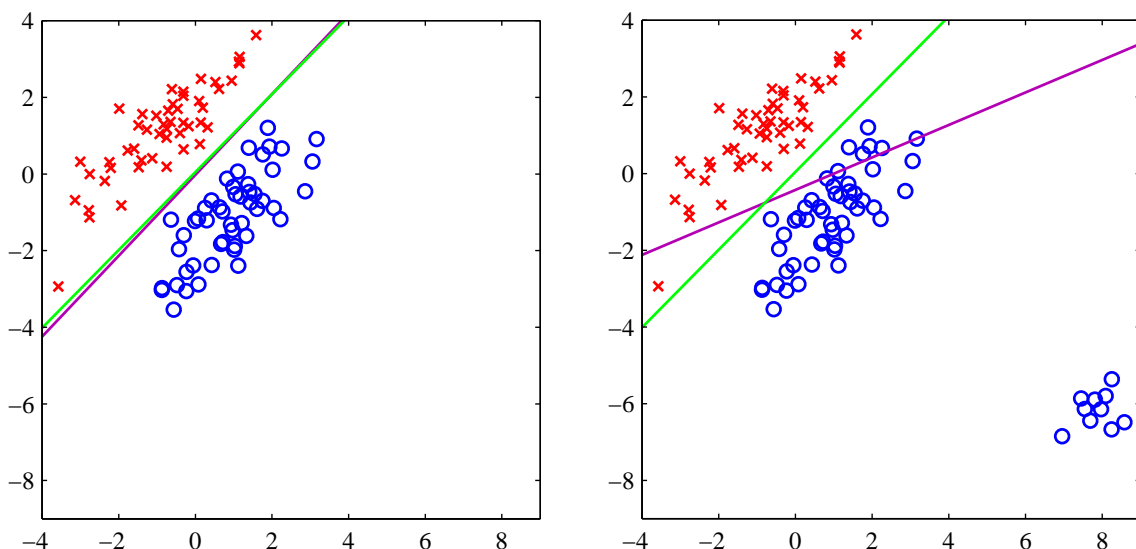
$$\mathbf{y}(\mathbf{x}) = \tilde{\mathbf{W}}^T \tilde{\mathbf{x}} = \begin{pmatrix} y_1(\mathbf{x}) \\ \vdots \\ y_K(\mathbf{x}) \end{pmatrix}$$

Assign class  $C_k$  to  $\mathbf{x}$ , where:

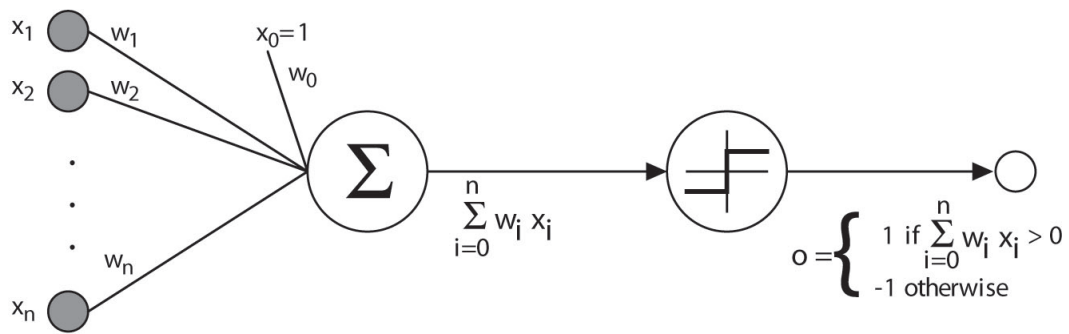
$$k = \underset{i \in \{1, \dots, K\}}{\operatorname{argmax}} \{y_i(\mathbf{x})\}$$

## Issues with least squares

Assume Gaussian conditional distributions. Not robust to outliers!



# Perceptron



$$o(x_1, \dots, x_d) = \begin{cases} 1 & \text{if } w_0 + w_1 x_1 + \dots + w_d x_d > 0 \\ -1 & \text{otherwise.} \end{cases}$$

$$o(\mathbf{x}) = \begin{cases} 1 & \text{if } \mathbf{w}^T \mathbf{x} > 0 \\ -1 & \text{otherwise.} \end{cases} = \text{sign}(\mathbf{w}^T \mathbf{x})$$

## Perceptron training rule

Consider the *unthresholded linear unit*, where

$$o = w_0 + w_1 x_1 + \dots + w_d x_d = \mathbf{w}^T \mathbf{x}$$

Let's learn  $w_i$  from training examples  $D = \{(\mathbf{x}_n, t_n)_{n=1}^N\}$  that minimize the squared error (*loss function*)

$$E(\mathbf{w}) \equiv \frac{1}{2} \sum_{n=1}^N (t_n - o_n)^2 = \frac{1}{2} \sum_{n=1}^N (t_n - \mathbf{w}^T \mathbf{x}_n)^2$$

## Perceptron training rule

$$\begin{aligned}
 \frac{\partial E}{\partial w_i} &= \frac{\partial}{\partial w_i} \frac{1}{2} \sum_{n=1}^N (t_n - \mathbf{w}^T \mathbf{x}_n)^2 = \frac{1}{2} \sum_{n=1}^N \frac{\partial}{\partial w_i} (t_n - \mathbf{w}^T \mathbf{x}_n)^2 \\
 &= \frac{1}{2} \sum_{n=1}^N 2(t_n - \mathbf{w}^T \mathbf{x}_n) \frac{\partial}{\partial w_i} (t_n - \mathbf{w}^T \mathbf{x}_n) \\
 &= \sum_{n=1}^N (t_n - \mathbf{w}^T \mathbf{x}_n) \frac{\partial}{\partial w_i} (t_n - \mathbf{w}^T \mathbf{x}_n) \\
 &= \sum_{n=1}^N (t_n - \mathbf{w}^T \mathbf{x}_n) (-x_{i,n})
 \end{aligned}$$

## Perceptron training rule

Unthresholded unit:

Update of weights  $\mathbf{w}$

$$\begin{aligned}
 w_i &\leftarrow w_i + \Delta w_i \\
 \Delta w_i &= -\eta \frac{\partial E}{\partial w_i} = \eta \sum_{n=1}^N (t_n - \mathbf{w}^T \mathbf{x}_n) x_{i,n}
 \end{aligned}$$

$\eta$  is a small constant (e.g., 0.05) called *learning rate*

## Perceptron training rule

Thresholded unit:

Update of weights  $\mathbf{w}$

$$w_i \leftarrow w_i + \Delta w_i$$

$$\Delta w_i = -\eta \frac{\partial E}{\partial w_i} = \eta \sum_{n=1}^N (t_n - \text{sign}(\mathbf{w}^T \mathbf{x}_n)) x_{i,n}$$

## Perceptron algorithm

Given perceptron model  $o(\mathbf{x}) = \text{sign}(\mathbf{w}^T \mathbf{x})$  and data set  $D$ , determine weights  $\mathbf{w}$ .

- 1 Initialize  $\hat{\mathbf{w}}$  (e.g. small random values)
- 2 Repeat until termination condition
  - $\hat{w}_i \leftarrow \hat{w}_i + \Delta w_i$
- 3 Output  $\hat{\mathbf{w}}$

# Perceptron algorithm

**Batch mode:** Consider all dataset  $D$

$$\Delta w_i = \eta \sum_{(\mathbf{x}, t) \in D} (t - o(\mathbf{x})) x_i$$

**Mini-Batch mode:** Choose a small subset  $S \subset D$

$$\Delta w_i = \eta \sum_{(\mathbf{x}, t) \in S} (t - o(\mathbf{x})) x_i$$

**Incremental mode:** Choose one sample  $(\mathbf{x}, t) \in D$

$$\Delta w_i = \eta (t - o(\mathbf{x})) x_i$$

$o(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$  for unthresholded,  $o(\mathbf{x}) = \text{sign}(\mathbf{w}^T \mathbf{x})$  for thresholded

Incremental and mini-batch modes speed up convergence and are less sensitive to local minima.

# Perceptron algorithm

Termination conditions

- Predefined number of iterations
- Threshold on changes in the loss function  $E(\mathbf{w})$

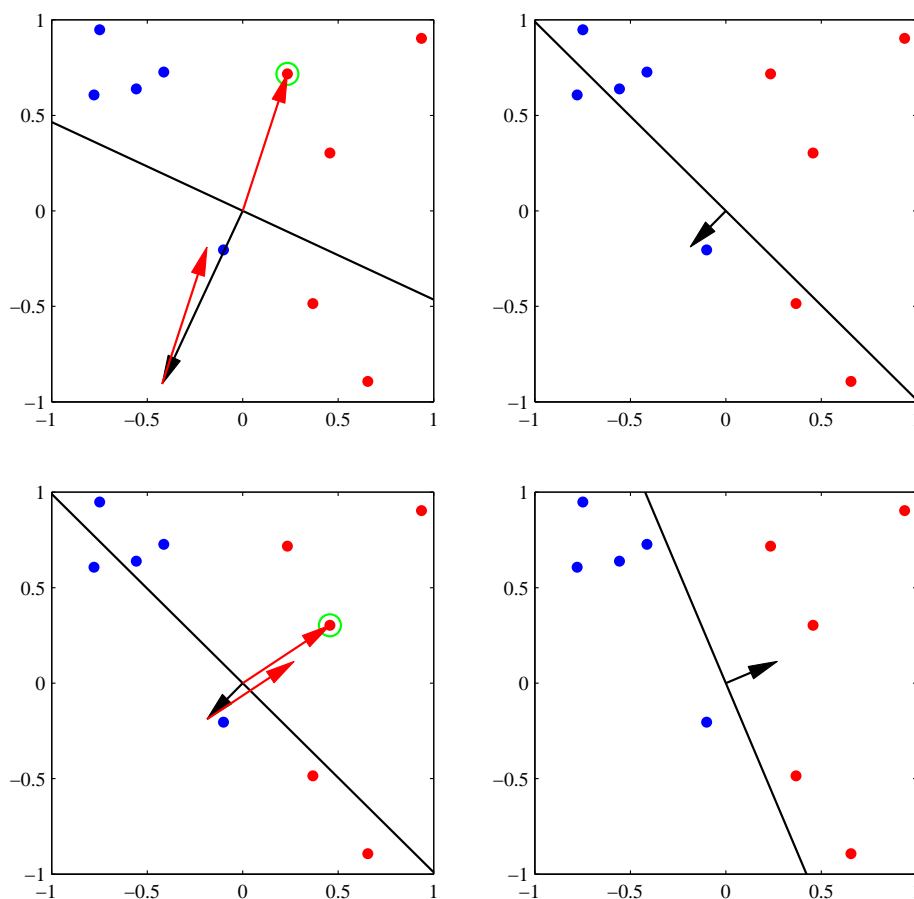
# Perceptron training rule

Example:

$$\eta = 0.1, x_i = 0.8$$

- if  $t = 1$  and  $o = -1$  then  $\Delta w_i = 0.16$
- if  $t = -1$  and  $o = 1$  then  $\Delta w_i = -0.16$

# Perceptron training rule



## Perceptron training rule

Can prove it will converge:

- if training data is linearly separable
- and  $\eta$  sufficiently small

Small  $\eta \rightarrow$  slow convergence.

## Perceptron: Prediction

Classification of new instance  $\mathbf{x}$  not in dataset:

Classify  $\mathbf{x}$  as  $C_k$ , for  $k = \text{sign}(\mathbf{w}^T \mathbf{x})$ , using learnt  $\mathbf{w}$

## Fisher's linear discriminant

Consider two classes case.

Determine  $y = \mathbf{w}^T \mathbf{x}$   
and classify  $\mathbf{x} \in C_1$  if  $y \geq -w_0$ ,  $\mathbf{x} \in C_2$  otherwise.

Corresponding to the projection on a line determined by  $\mathbf{w}$ .

## Fisher's linear discriminant

Adjusting  $\mathbf{w}$  to find a direction that maximizes class separation.

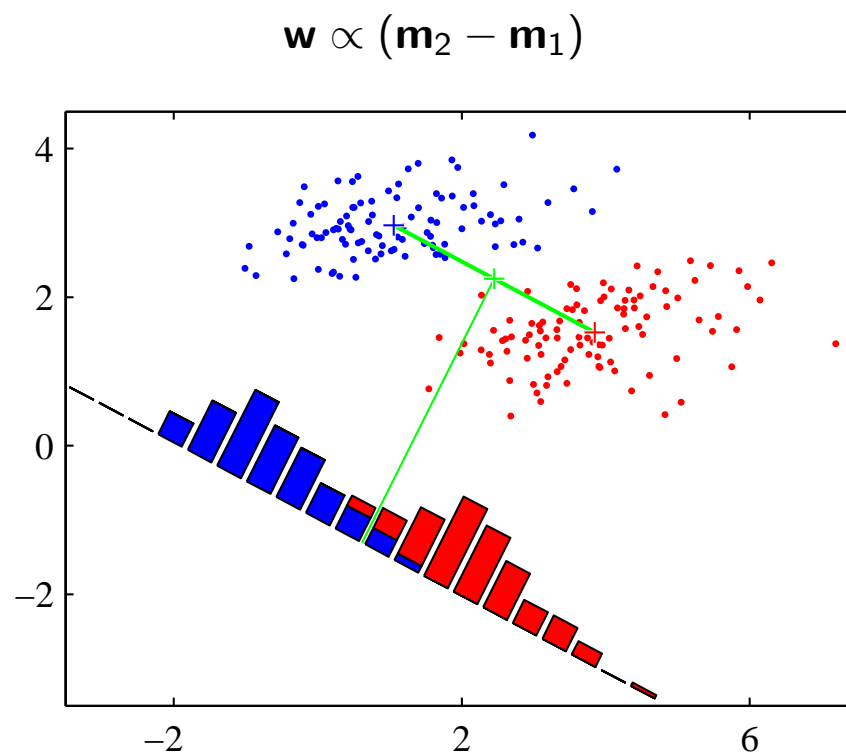
Consider a data set with  $N_1$  points in  $C_1$  and  $N_2$  points in  $C_2$

$$\mathbf{m}_1 = \frac{1}{N_1} \sum_{n \in C_1} \mathbf{x}_n \quad \mathbf{m}_2 = \frac{1}{N_2} \sum_{n \in C_2} \mathbf{x}_n$$

Choose  $\mathbf{w}$  that maximizes  $J(\mathbf{w}) = \mathbf{w}^T (\mathbf{m}_2 - \mathbf{m}_1)$ , subject to  $\|\mathbf{w}\| = 1$ .



# Fisher's linear discriminant



# Fisher's linear discriminant

Fisher criterion

$$J(\mathbf{w}) = \frac{\mathbf{w}^T \mathbf{S}_B \mathbf{w}}{\mathbf{w}^T \mathbf{S}_W \mathbf{w}}$$

with

$$\mathbf{S}_B = (\mathbf{m}_2 - \mathbf{m}_1)(\mathbf{m}_2 - \mathbf{m}_1)^T$$

Between class scatter

$$\mathbf{S}_W = \sum_{n \in C_1} (\mathbf{x}_n - \mathbf{m}_1)(\mathbf{x}_n - \mathbf{m}_1)^T + \sum_{n \in C_2} (\mathbf{x}_n - \mathbf{m}_2)(\mathbf{x}_n - \mathbf{m}_2)^T$$

Within class scatter

Choose  $\mathbf{w}$  that maximizes  $J(\mathbf{w})$ .

# Fisher's linear discriminant

Find  $\mathbf{w}$  that maximizes

$$J(\mathbf{w}) = \frac{\mathbf{w}^T \mathbf{S}_B \mathbf{w}}{\mathbf{w}^T \mathbf{S}_W \mathbf{w}}$$

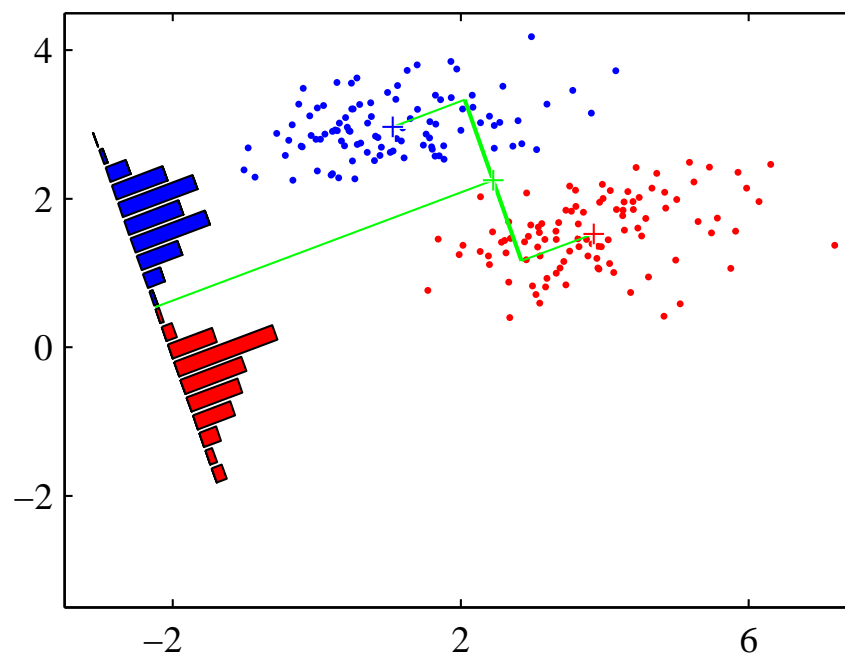
by solving

$$\frac{d}{d\mathbf{w}} J(\mathbf{w}) = 0$$

$$\Rightarrow \mathbf{w}^* \propto \mathbf{S}_W^{-1}(\mathbf{m}_2 - \mathbf{m}_1)$$

# Fisher's linear discriminant

$$\mathbf{w} \propto \mathbf{S}_W^{-1}(\mathbf{m}_2 - \mathbf{m}_1)$$



## Fisher's linear discriminant

Summarizing, given a two classes classification problem, Fisher's linear discriminant is given by the function  $y = \mathbf{w}^T \mathbf{x}$  and the classification of new instances is given by  $y \geq -w_0$  where

$$\mathbf{w} = \mathbf{S}_W^{-1}(\mathbf{m}_2 - \mathbf{m}_1)$$

$$w_0 = \mathbf{w}^T \mathbf{m}$$

$\mathbf{m}$  is the global mean of all the data set.

## Fisher's linear discriminant

Multiple classes.

$$\mathbf{y} = \mathbf{W}^T \mathbf{x}$$

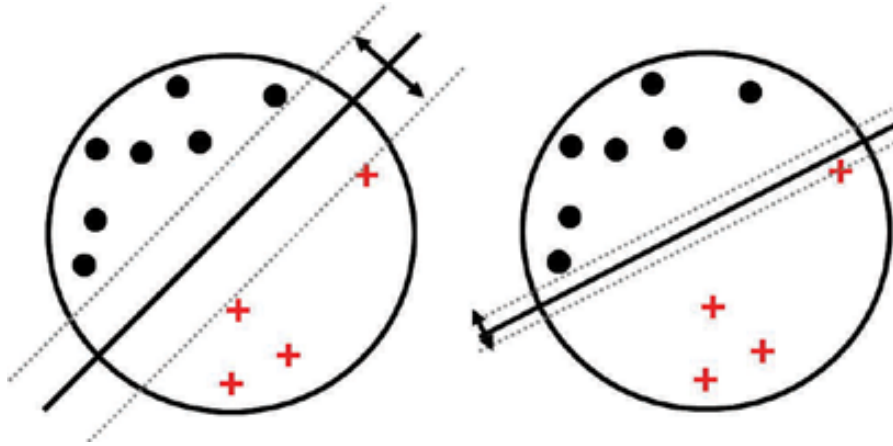
Maximizing

$$J(\mathbf{W}) = \text{Tr} \left\{ (\mathbf{W} \mathbf{S}_W \mathbf{W}^T)^{-1} (\mathbf{W} \mathbf{S}_B \mathbf{W}^T) \right\}$$

...

# Support Vector Machines

Support Vector Machines (SVM) for Classification aims at maximum margin providing for better accuracy.



# Support Vector Machines

Let's consider binary classification  $y : X \rightarrow \{+1, -1\}$  with data set  $D = \{(\mathbf{x}_n, t_n)_{n=1}^N\}$ ,  $t_n \in \{+1, -1\}$  and a linear model

$$y(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0$$

Assume  $D$  is linearly separable

$$\exists \mathbf{w}, w_0 \text{ s.t. } \begin{cases} y(\mathbf{x}_n) > 0, & \text{if } t_n = +1 \\ y(\mathbf{x}_n) < 0, & \text{if } t_n = -1 \end{cases}$$

$$t_n y(\mathbf{x}_n) > 0 \quad \forall n = 1, \dots, N$$

# Support Vector Machines

Let  $\mathbf{x}_k$  be the closest point of the data set  $D$  to the hyperplane  $\bar{h} : \bar{\mathbf{w}}^T \mathbf{x} + \bar{w}_0 = 0$

the *margin* (smallest distance between  $\mathbf{x}_k$  and  $\bar{h}$ ) is  $\frac{|y(\mathbf{x}_k)|}{\|\bar{\mathbf{w}}\|}$

Given data set  $D$  and hyperplane  $\bar{h}$ , the margin is computed as

$$\min_{n=1,\dots,N} \frac{|y(\mathbf{x}_n)|}{\|\bar{\mathbf{w}}\|} = \dots = \frac{1}{\|\bar{\mathbf{w}}\|} \min_{n=1,\dots,N} [t_n(\bar{\mathbf{w}}^T \mathbf{x}_n + \bar{w}_0)]$$

using the property  $|y(\mathbf{x}_n)| = t_n y(\mathbf{x}_n)$

# Support Vector Machines

Given data set  $D$ , the hyperplane  $h^* : \mathbf{w}^{*T} \mathbf{x} + w_0^* = 0$  with maximum margin is computed as

$$\mathbf{w}^*, w_0^* = \operatorname{argmax}_{\mathbf{w}, w_0} \frac{1}{\|\mathbf{w}\|} \min_{n=1,\dots,N} [t_n(\mathbf{w}^T \mathbf{x}_n + w_0)]$$

# Support Vector Machines

Rescaling all the points does not affect the solution.

Rescale in such a way that for the closet point  $\mathbf{x}_k$  we have

$$t_k(\mathbf{w}^T \mathbf{x}_k + w_0) = 1$$

Canonical representation:

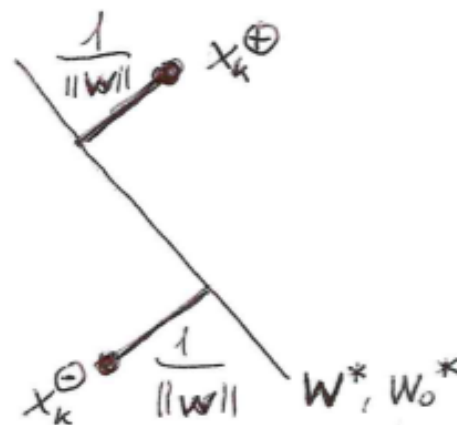
$$t_n(\mathbf{w}^T \mathbf{x}_n + w_0) \geq 1 \quad \forall n = 1, \dots, N$$

# Support Vector Machines

When the maximum margin hyperplane  $\mathbf{w}^*, w_0^*$  is found, there will be at least 2 closest points  $\mathbf{x}_k^+$  and  $\mathbf{x}_k^-$  (one for each class).

$$\mathbf{w}^{*T} \mathbf{x}_k^+ + w_0^* = +1$$

$$\mathbf{w}^{*T} \mathbf{x}_k^- + w_0^* = -1$$



# Support Vector Machines

In the canonical representation of the problem the maximum margin hyperplane can be found by solving the optimization problem

$$\mathbf{w}^*, w_0^* = \operatorname{argmax} \frac{1}{\|\mathbf{w}\|} = \operatorname{argmin} \frac{1}{2} \|\mathbf{w}\|^2$$

subject to

$$t_n(\mathbf{w}^T \mathbf{x}_n + w_0) \geq 1 \quad \forall n = 1, \dots, N$$

Quadratic programming problem solved with Lagrangian method.

# Support Vector Machines

Solution

$$\mathbf{w}^* = \sum_{n=1}^N a_n^* t_n \mathbf{x}_n$$

$a_i^*$  (Lagrange multipliers): results of the Lagrangian optimization problem

$$\tilde{L}(\mathbf{a}) = \sum_{n=1}^N a_n - \frac{1}{2} \sum_{n=1}^N \sum_{m=1}^N a_n a_m t_n t_m \mathbf{x}_n^T \mathbf{x}_m$$

subject to

$$a_n \geq 0 \quad \forall n = 1, \dots, N$$

$$\sum_{n=1}^N a_n t_n = 0$$

# Support Vector Machines

Note:

samples  $\mathbf{x}_n$  for which  $a_n^* = 0$  do not contribute to the solution

Karush-Kuhn-Tucker (KKT) condition:

for each  $\mathbf{x}_n \in D$ , either  $a_n^* = 0$  or  $t_n y(\mathbf{x}_n) = 1$

thus  $t_n y(\mathbf{x}_n) > 1$  implies  $a_n^* = 0$

Support vectors:  $\mathbf{x}_k$  such that  $t_k y(\mathbf{x}_k) = 1$  and  $a_k^* > 0$

$$SV \equiv \{\mathbf{x}_k \in D \mid t_k y(\mathbf{x}_k) = 1\}$$

# Support Vector Machines

Hyperplanes expressed with support vectors

$$y(\mathbf{x}) = \sum_{\mathbf{x}_j \in SV} a_j^* t_j \mathbf{x}^T \mathbf{x}_j + w_0^* = 0$$

Note: other vectors  $\mathbf{x}_n \notin SV$  do not contribute ( $a_n^* = 0$ )



# Support Vector Machines

To compute  $w_0^*$ :

Support vector  $\mathbf{x}_k \in SV$  satisfies  $t_k y(\mathbf{x}_k) = 1$

$$t_k \left( \sum_{\mathbf{x}_j \in SV} a_j^* t_j \mathbf{x}_k^T \mathbf{x}_j + w_0^* \right) = 1$$

Multiplying by  $t_k$  and using  $t_k^2 = 1$

$$w_0^* = t_k - \sum_{\mathbf{x}_j \in SV} a_j^* t_j \mathbf{x}_k^T \mathbf{x}_j$$

# Support Vector Machines

Instead of using one particular support vector  $\mathbf{x}_k$  to determine  $w_0$

$$w_0^* = t_k - \sum_{\mathbf{x}_j \in SV} a_j^* t_j \mathbf{x}_k^T \mathbf{x}_j$$

a more stable solution is obtained by averaging over all the support vectors

$$w_0^* = \frac{1}{|SV|} \sum_{\mathbf{x}_k \in SV} \left( t_k - \sum_{\mathbf{x}_j \in S} a_j^* t_j \mathbf{x}_k^T \mathbf{x}_j \right)$$

# Support Vector Machines

Given the maximum margin hyperplane determined by  $a_k^*$ ,  $w_0^*$

Classification of a new instance  $\mathbf{x}'$

$$y(\mathbf{x}') = \text{sign} \left( \sum_{\mathbf{x}_k \in SV} a_k^* t_k \mathbf{x}'^T \mathbf{x}_k + w_0^* \right)$$

# Support Vector Machines

Optimization problem for determining  $\mathbf{w}$ ,  $w_0$  (dimension  $d + 1$ , with  $X = \mathbb{R}^d$ ) transformed in an optimization problem for determining  $\mathbf{a}$  (dimension  $|D|$ )

Efficient when  $d \ll |D|$  (most of  $a_i$  will be zero).

Very useful when  $d$  is large or infinite.


# Support Vector Machines with soft margin constraints

What if data are “almost” linearly separable (e.g., a few points are on the “wrong side”)

Let us introduce *slack variables*  $\xi_n \geq 0$   $n = 1, \dots, N$

# Support Vector Machines with soft margin constraints

- $\xi_n = 0$  if point on or inside the correct margin boundary
- $0 < \xi_n \leq 1$  if point inside the margin but correct side
- $\xi_n > 1$  if point on wrong side of boundary



when  $\xi_n = 1$ , the sample lies on the decision boundary  $y(\mathbf{x}_n) = 0$   
 when  $\xi_n > 1$ , the sample will be mis-classified

# Support Vector Machines with soft margin constraints

*Soft margin* constraint

$$t_n y(\mathbf{x}_n) \geq 1 - \xi_n, \quad n = 1, \dots, N$$

Optimization problem with soft margin constraints

$$\mathbf{w}^*, w_0^* = \operatorname{argmin} \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{n=1}^N \xi_n$$

subject to

$$t_n y(\mathbf{x}_n) \geq 1 - \xi_n, \quad n = 1, \dots, N$$

$$\xi_n \geq 0, \quad n = 1, \dots, N$$

$C$  is a constant (inverse of a regularization coefficient)

# Support Vector Machines with soft margin constraints

Solution similar to the case of linearly separable data.

$$\mathbf{w}^* = \sum_{n=1}^N a_n^* t_n \mathbf{x}_n$$

$$w_0^* = \dots$$

with  $a_n^*$  computed as solution of a Lagrangian optimization problem.

# Basis functions

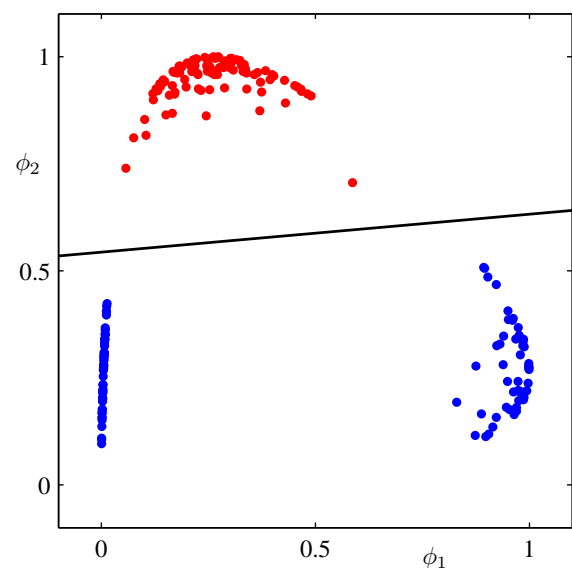
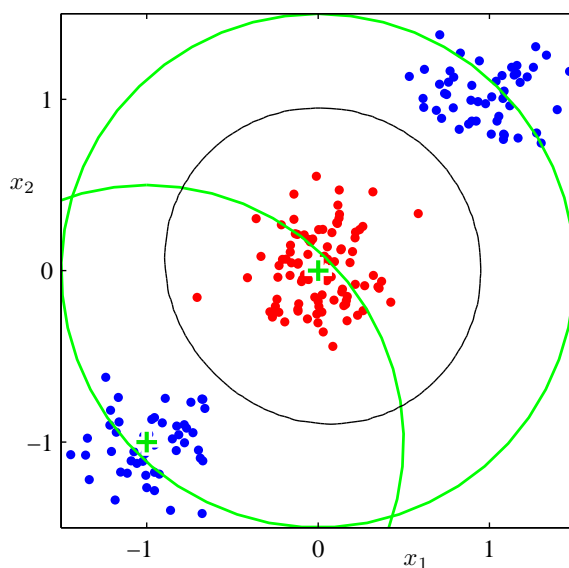
So far we considered models working directly on  $\mathbf{x}$ .

All the results hold if we consider a non-linear transformation of the inputs  $\phi(\mathbf{x})$  (*basis functions*).

Decision boundaries will be linear in the feature space  $\phi$  and non-linear in the original space  $\mathbf{x}$

Classes that are linearly separable in the feature space  $\phi$  may not be separable in the input space  $\mathbf{x}$ .

## Basis functions example



## Basis functions examples

- Linear
- Polynomial
- Radial Basis Function (RBF)
- Sigmoid
- ...

## Linear models for non-linear functions

Learning non-linear function

$$y : X \rightarrow \{C_1, \dots, C_K\}$$

from data set  $D$  non-linearly separable.

Find a non-linear transformation  $\phi$  and learn a linear model

$$y(\mathbf{x}) = \mathbf{w}^T \phi(\mathbf{x}) + w_0 \text{ (two classes)}$$

$$y_k(\mathbf{x}) = \mathbf{w}_k^T \phi(\mathbf{x}) + w_{k0} \text{ (multiple classes)}$$

# Summary

- Basic methods for learning linear classification functions
- Based on solution of an optimization problem
- Closed form vs. iterative solutions
- Sensitivity to outliers
- Learning non-linear functions with linear models using basis functions
- Further developed as kernel methods